

Mid-Level Engineers Persistent Candidates - Payments Take-Home Problems

This exercise evaluates how you think, structure, and solve problems, not just your ability to code a perfect solution. We're looking for clean, readable code, reasonable architecture, and thoughtful tradeoffs. We're not looking for perfection — we want to understand how you work and think.

Tools Required [🔗](#)

You are expected to complete this exercise using [Cursor](#), an AI-assisted IDE designed for developers.

- You can start with a free trial of Cursor if you don't already have an account.
- The goal is not just to write code, but to show how you collaborate with an AI development tool.

What We're Evaluating [🔗](#)

- Clarity and maintainability of your code
- Approach to problem solving and handling edge cases
- Use of modern development practices (e.g., modularity, tests, API design)
- Your ability to work independently, yet communicate decisions clearly
- How effectively you use AI (via Cursor) to aid development
- Creativity in applying stretch goals or AI features (for senior-level roles)

Submission Guidelines [🔗](#)

- Share a public GitHub repo or zipped file
- Include a `README.md` with:
 - Setup instructions
 - Description of fraud logic and LLM usage
 - Explanation of any assumptions or tradeoffs

How We Will Review Your Submission [🔗](#)

We will:

- Run your project locally (or using Docker, if provided)
- Review your codebase structure and naming conventions
- Review your Cursor prompts history
- Discuss your approach in a follow-up interview
- Ask how you used AI to help you think through problems, debug, or generate code

Mid-Level Problem 1: Mini Payment Gateway Proxy with LLM Risk Summary [🔗](#)

Design a lightweight backend API that simulates routing payment requests to Stripe or PayPal based on a simple fraud risk score. Use an LLM to generate a human-readable explanation of each decision.

Requirements [🔗](#)

1. API Endpoint: [🔗](#)

```
1 POST /charge
```

Request Payload: [🔗](#)

```
1 {
2   "amount": 1000,
3   "currency": "USD",
4   "source": "tok_test",
5   "email": "donor@example.com"
6 }
```

Behavior: [🔗](#)

- Internally simulate a fraud risk score (float between 0 and 1)
 - Use simple heuristics: large amount, suspicious domain (e.g., `.ru`, `test.com`)
- Based on the score:
 - Route to payment processor if score < 0.5
 - Block submission if score ≥ 0.5
- Generate a natural-language summary of the risk using an LLM

Example Response [🔗](#)

```
1 {
2   "transactionId": "txn_abc123",
3   "provider": "paypal",
4   "status": "success",
5   "riskScore": 0.32,
6   "explanation": "This payment was routed to PayPal due to a moderately high low score based on a large amount
and a suspicious email domain."
7 }
```

- Log all transactions in memory with timestamp and metadata
- Handle input validation and basic error responses
- Include basic unit tests
- Use modern TypeScript conventions and structure

Stretch Goals [🔗](#)

- Add a `GET /transactions` endpoint
- Cache frequent LLM prompts for performance
- Dockerize the app
- Add a basic fraud rule config file to make heuristics configurable

Mid-Level Problem 2: Subscription Billing Simulator [🔗](#)

Build a backend service that simulates recurring donations for nonprofit supporters. You'll handle subscription creation, automatic charging based on schedule, and use an LLM to tag and summarize each campaign description.

Requirements [🔗](#)

1. API Endpoints [🔗](#)

- `POST /subscriptions`

```
1 {
2   "donorId": "abc123",
3   "amount": 1500,
```

```
4   "currency": "USD",
5   "interval": "monthly",
6   "campaignDescription": "Emergency food and clean water for earthquake victims in Nepal"
7 }
```

- Store each active subscription
- Call an LLM to:
 - Tag the campaign (e.g., "disaster relief", "Nepal")
 - Summarize it in one sentence
 - Store these as part of the subscription record
- DELETE /subscriptions/:donorId
 - Cancels a recurring donation

2. Daily/Batched Processing [🔗](#)

- Simulate a background job (e.g., setInterval) that:
 - “Charges” active subscriptions based on their interval
 - Logs a new transaction for each recurring payment
 - Output can be a console log or stored in memory

3. LLM Integration [🔗](#)

Use a LLM to generate:

- Tags from campaignDescription
- Short summary of the campaign

Example LLM Output:

```
1 {
2   "tags": ["disaster relief", "clean water", "Nepal"],
3   "summary": "This campaign provides emergency aid to earthquake-affected communities in Nepal."
4 }
```

Stretch Goals [🔗](#)

- Add a GET /subscriptions endpoint to list all active subscriptions
- Add a GET /transactions endpoint to view donation history
- Add support for multiple currencies and normalize amounts (e.g., in USD)
- Containerize using Docker